

Final Project Submission

Please fill out: Student names: Trixie Cherop Josephine Wanjiru Evalyne Macharia Mercy Cherotich Priscillah Veke Laurah Mutheu Student pace: part time Scheduled project review date/time: Instructor name: Blog post URL:

ANALYSIS OF KEY INDICATORS OF HOUSE PRICES

Research Objectives

Main Objective

To build a linear Regression Model that predicts House Prices

Specific Objectives

To Identify key features that influence house House prices

To assess the feature with the highest impact on House prices

To evaluate and validate the performance of the model

Business Problem

Real estate is a highly dynamic market influenced by numerous factors. This makes it challenging for real estate investors to accurately predict house prices. Inaccurate pricing models can lead to reduced profitability, missed opportunities, and dissatisfied customers. The current pricing strategy of the real estate company is suboptimal, leading to potential loss of revenue and increased customer dissatisfaction. Hence, the need of a robust predictive pricing model to enable companies stay competitive and adapt to market fluctuations.

Key Challenges:

Difficulty in identifying the most influential features impacting house prices.

Inability to accurately predict house prices based on relevant features.

Limited understanding of the factors driving property value in the current market.

Lack of a data driven pricing strategy, leading to potential underpricing or overpricing of

Project Overview

This project is aimed at helping real estate investors make informed decision on what type of houses they should invest in. This is in terms of the most impactful features, both positively and negatively, on House prices. The key components of the analysis include Data preparation, Feature selection and Engineering, Model Development, Evaluation and Validation.

In [2]: ▶

```
# import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

# Loading the dataset
data=pd.read_csv('kc_house_data.csv')
data.head()
```

Out[2]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

◀

▶

In [3]: `data.tail()`

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0

5 rows × 21 columns

In [4]: `# checking summary`
`data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

Data preprocessing

Data cleaning

```
In [5]: ▶ # checking null values
null= data.isna().sum()
null
```

```
Out[5]: id                0
        date              0
        price             0
        bedrooms          0
        bathrooms         0
        sqft_living        0
        sqft_lot           0
        floors             0
        waterfront        2376
        view              63
        condition         0
        grade             0
        sqft_above         0
        sqft_basement      0
        yr_built           0
        yr_renovated       3842
        zipcode           0
        lat               0
        long              0
        sqft_living15      0
        sqft_lot15         0
        dtype: int64
```

```
In [6]: # percentage of missing data  
percentage_missing=null*100/len(data)  
percentage_missing
```

```
Out[6]: id            0.000000  
date            0.000000  
price           0.000000  
bedrooms        0.000000  
bathrooms       0.000000  
sqft_living     0.000000  
sqft_lot        0.000000  
floors          0.000000  
waterfront     11.001528  
view            0.291707  
condition       0.000000  
grade           0.000000  
sqft_above      0.000000  
sqft_basement   0.000000  
yr_built        0.000000  
yr_renovated    17.789508  
zipcode         0.000000  
lat             0.000000  
long            0.000000  
sqft_living15   0.000000  
sqft_lot15      0.000000  
dtype: float64
```

From the results above one of the variables for our analysis 'view' has some missing data of 0.291707%. We will proceed and first clean that.

```
In [7]: data["view"].unique()
```

```
Out[7]: array(['NONE', nan, 'GOOD', 'EXCELLENT', 'AVERAGE', 'FAIR'], dtype=object)
```

```
In [8]: # dealing with missing data on 'view' column  
# drop the null values for 'view' since it is a small percentage  
data.dropna(axis=0, subset=['view'], inplace=True)  
data["view"].isnull().sum()
```

```
Out[8]: 0
```

```
In [9]: # replace null values in column 'waterfront' with place holder 'unknown'  
data['waterfront'].fillna('Unknown', inplace=True)  
data["waterfront"].isnull().sum()
```

```
Out[9]: 0
```

```
In [10]: data["yr_renovated"].unique()
```

```
Out[10]: array([ 0., 1991., nan, 2002., 2010., 1992., 2013., 1994., 1978.,
        2005., 2003., 1984., 1954., 2014., 2011., 1983., 1945., 1990.,
        1988., 1977., 1981., 1995., 2000., 1999., 1998., 1970., 1989.,
        2004., 1986., 2007., 1987., 2006., 1985., 2001., 1980., 1971.,
        1979., 1997., 1950., 1969., 1948., 2009., 2015., 1974., 2008.,
        1968., 2012., 1963., 1951., 1962., 1953., 1993., 1996., 1955.,
        1982., 1956., 1940., 1976., 1946., 1975., 1964., 1973., 1957.,
        1959., 1960., 1967., 1965., 1934., 1972., 1944., 1958.]
```

```
In [11]: # replace null values in column with place holder '0'
data['yr_renovated'].fillna('0', inplace=True)
data["yr_renovated"].isnull().sum()
```

```
Out[11]: 0
```

```
In [12]: # checking if all missing data have been cleaned
data.isnull().sum()
```

```
Out[12]: id          0
        date         0
        price        0
        bedrooms     0
        bathrooms    0
        sqft_living   0
        sqft_lot      0
        floors        0
        waterfront    0
        view          0
        condition     0
        grade         0
        sqft_above    0
        sqft_basement 0
        yr_built      0
        yr_renovated  0
        zipcode       0
        lat           0
        long          0
        sqft_living15 0
        sqft_lot15    0
        dtype: int64
```

We see that all the missing values have been cleaned

Dealing with categorical variables

One-hot encoding

We are going to encode the categorical variables, 'grade', 'view', 'waterfront', 'condition' to numeric

```
In [13]: #encoding 'grade' column
data['grade'].unique()
```

```
Out[13]: array(['7 Average', '6 Low Average', '8 Good', '11 Excellent', '9 Better',
               '5 Fair', '10 Very Good', '12 Luxury', '4 Low', '3 Poor',
               '13 Mansion'], dtype=object)
```

```
In [14]: # getting dummy variables
dummy_grade = pd.get_dummies(data['grade'], prefix='grade')

# Concatenate the dummy variables with the original DataFrame
data = pd.concat([data, dummy_grade], axis=1)

# Dropping the original 'grade' column
data = data.drop('grade', axis=1)
data = data.replace({True: 1, False: 0})
```

```
In [15]: data.head()
```

```
Out[15]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	1
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	0
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	0
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	0
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	0

5 rows × 31 columns



```
In [16]: #encoding 'view' column
data['view'].unique()
```

```
Out[16]: array(['NONE', 'GOOD', 'EXCELLENT', 'AVERAGE', 'FAIR'], dtype=object)
```

```
In [17]: # getting dummies
dummy_view = pd.get_dummies(data['view'], prefix='view')

#Concatenate the dummy variables with the original DataFrame
data = pd.concat([data, dummy_view], axis=1)


# Dropping the original 'view' column
data = data.drop('view', axis=1)
data = data.replace({True: 1, False: 0})
```

```
In [18]: data.head()
```

Out[18]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	l
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 35 columns



```
In [19]: #encoding 'waterfront' column
data['waterfront'].unique()
```

Out[19]: array(['Unknown', 'NO', 'YES'], dtype=object)

```
In [20]: # getting dummies
dummy_waterfront = pd.get_dummies(data['waterfront'], prefix='waterfront')

#Concatenate the dummy variables with the original DataFrame
data = pd.concat([data, dummy_waterfront], axis=1)

# Dropping the original 'condition' column
data = data.drop('waterfront', axis=1)


data = data.replace({True: 1, False: 0})
```


In [21]: `data.head()`

Out[21]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	co
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	/
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	/
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	/
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	/

5 rows × 37 columns



In [22]: `#encoding 'condition' column
data['condition'].unique()`

Out[22]: `array(['Average', 'Very Good', 'Good', 'Poor', 'Fair'], dtype=object)`

In [23]: `# getting dummies
dummy_condition = pd.get_dummies(data['condition'], prefix='condition')

#Concatenate the dummy variables with the original DataFrame
data = pd.concat([data, dummy_condition], axis=1)


Dropping the original 'condition' column
data = data.drop('condition', axis=1)
data = data.replace({True: 1, False: 0})`

In [24]: `data.head()`

Out[24]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sq
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 41 columns



```
In [25]: ▶ data['sqft_basement'].unique()
```



```
'946.0', '1281.0', '414.0', '276.0', '1248.0', '602.0', '516.0',  
'176.0', '225.0', '1275.0', '266.0', '283.0', '65.0', '2310.0',  
'10.0', '1770.0', '2120.0', '295.0', '207.0', '915.0', '556.0',  
'417.0', '143.0', '508.0', '2810.0', '20.0', '274.0', '248.0'],  
dtype=object)
```

```
In [26]: ▶ dropping_question_mark = data[data['sqft_basement'] == '?']  
data = data.drop(dropping_question_mark.index )
```

```
In [27]: ▶ # changing data type of 'sqft_basement' to float  
data['sqft_basement'] = data['sqft_basement'].astype('float64')
```

In [28]: `data.dtypes`

```
Out[28]: id                int64
date                object
price              float64
bedrooms           int64
bathrooms          float64
sqft_living        int64
sqft_lot           int64
floors             float64
sqft_above         int64
sqft_basement      float64
yr_built           int64
yr_renovated       object
zipcode            int64
lat               float64
long              float64
sqft_living15      int64
sqft_lot15         int64
grade_10 Very Good int64
grade_11 Excellent int64
grade_12 Luxury    int64
grade_13 Mansion   int64
grade_3 Poor       int64
grade_4 Low        int64
grade_5 Fair       int64
grade_6 Low Average int64
grade_7 Average    int64
grade_8 Good       int64
grade_9 Better     int64
view_AVERAGE      int64
view_EXCELLENT     int64
view_FAIR          int64
view_GOOD          int64
view_NONE          int64
waterfront_NO      int64
waterfront_Unknown int64
waterfront_YES     int64
condition_Average  int64
condition_Fair     int64
condition_Good     int64
condition_Poor     int64
condition_Very Good int64
dtype: object
```

Exploratory Data Analysis

In [29]: `# checking rows and columns`
`data.shape`

```
Out[29]: (21082, 41)
```

```
In [30]:  # checking data types  
data.dtypes
```

```
Out[30]: id                int64  
date                object  
price               float64  
bedrooms            int64  
bathrooms           float64  
sqft_living          int64  
sqft_lot             int64  
floors              float64  
sqft_above           int64  
sqft_basement        float64  
yr_built             int64  
yr_renovated         object  
zipcode             int64  
lat                 float64  
long                float64  
sqft_living15        int64  
sqft_lot15           int64  
grade_10 Very Good   int64  
grade_11 Excellent   int64  
grade_12 Luxury       int64  
grade_13 Mansion     int64  
grade_3 Poor          int64  
grade_4 Low           int64  
grade_5 Fair          int64  
grade_6 Low Average   int64  
grade_7 Average       int64  
grade_8 Good          int64  
grade_9 Better        int64  
view_AVERAGE         int64  
view_EXCELLENT        int64  
view_FAIR             int64  
view_GOOD             int64  
view_NONE             int64  
waterfront_NO         int64  
waterfront_Unknown    int64  
waterfront_YES        int64  
condition_Average     int64  
condition_Fair         int64  
condition_Good         int64  
condition_Poor         int64  
condition_Very Good   int64  
dtype: object
```

```
In [31]: ▶ # checking columns
data.columns
```

```
Out[31]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
               'sqft_lot', 'floors', 'sqft_above', 'sqft_basement', 'yr_built',
               'yr_renovated', 'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_
               lot15',
               'grade_10 Very Good', 'grade_11 Excellent', 'grade_12 Luxury',
               'grade_13 Mansion', 'grade_3 Poor', 'grade_4 Low', 'grade_5 Fai
               r',
               'grade_6 Low Average', 'grade_7 Average', 'grade_8 Good',
               'grade_9 Better', 'view_AVERAGE', 'view_EXCELLENT', 'view_FAIR',
               'view_GOOD', 'view_NONE', 'waterfront_NO', 'waterfront_Unknown',
               'waterfront_YES', 'condition_Average', 'condition_Fair',
               'condition_Good', 'condition_Poor', 'condition_Very Good'],
               dtype='object')
```

```
In [32]: ▶ # dropping columns

data = data.drop(['id', 'date', 'yr_renovated'], axis=1)
```

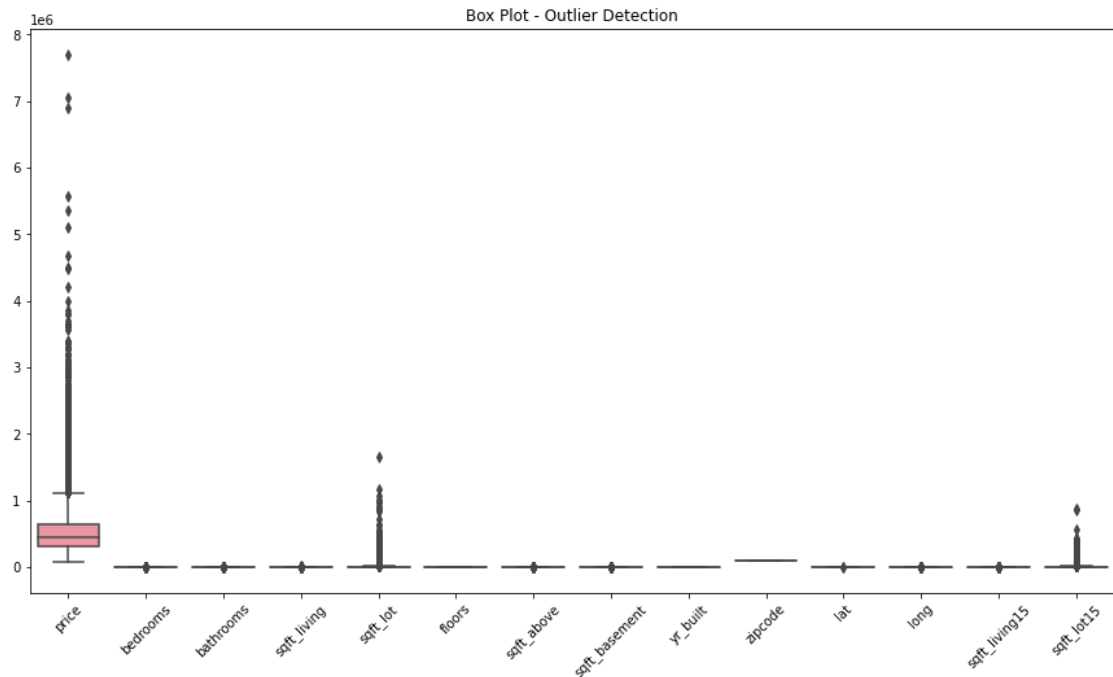
Checking outliers

```
In [33]: ▶ # Create box plots to visualize outliers
plt.figure(figsize=(15, 8))
sns.boxplot(data=data[['price', 'bedrooms', 'bathrooms', 'sqft_living',
                        'sqft_lot', 'floors', 'sqft_above', 'sqft_basement', 'yr_built',
                        'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15']])
plt.title('Box Plot - Outlier Detection')
plt.xticks(rotation=45)
plt.show()

# Calculating z-scores for numerical features
numeric_features = ['price', 'bedrooms', 'bathrooms', 'sqft_living',
                    'sqft_lot', 'floors', 'sqft_above', 'sqft_basement', 'yr_built',
                    'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15']
z_scores = data[numeric_features].apply(lambda x: (x - x.mean()) / x.std())

# Identify outliers based on z-score threshold ( z-score > 3 or z-score <
outliers = data[(z_scores > 3).any(axis=1)]

# Print the outliers
print('Outliers:')
print(outliers)
```



Outliers:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
5	1230000.0	4	4.50	5420	101930	1.0	
10	662500.0	3	2.50	3560	9796	1.0	
21	2000000.0	3	2.75	3050	44867	1.0	
41	775000.0	4	2.25	4220	24186	1.0	
70	1040000.0	5	3.25	4770	50094	1.0	
...	
21545	750000.0	5	4.00	4500	8130	2.0	
21552	1700000.0	4	3.50	3830	8963	2.0	
21560	3570000.0	5	4.50	4850	10584	2.0	
21574	1220000.0	4	3.50	4910	9444	1.5	
21584	1540000.0	5	3.75	4470	8088	2.0	

	sqft_above	sqft_basement	yr_built	zipcode	...	view_GOOD	\
5	3890	1530.0	2001	98053	...	0	
10	1860	1700.0	1965	98007	...	0	
21	2330	720.0	1968	98040	...	0	
41	2600	1620.0	1984	98166	...	0	
70	3070	1700.0	1973	98005	...	0	
...	
21545	4500	0.0	2007	98059	...	0	
21552	3120	710.0	2014	98004	...	0	
21560	3540	1310.0	2007	98008	...	0	
21574	3110	1800.0	2007	98074	...	0	
21584	4470	0.0	2008	98004	...	0	

	view_NONE	waterfront_NO	waterfront_Unknown	waterfront_YES	\
5	1	1	0	0	
10	1	0	1	0	
21	0	1	0	0	
41	1	1	0	0	
70	1	1	0	0	
...	
21545	1	0	1	0	
21552	1	1	0	0	
21560	0	0	0	1	
21574	1	1	0	0	
21584	1	1	0	0	

	condition_Average	condition_Fair	condition_Good	condition_Poor	\
5	1	0	0	0	
10	1	0	0	0	
21	1	0	0	0	
41	1	0	0	0	
70	0	0	1	0	
...	
21545	1	0	0	0	
21552	1	0	0	0	
21560	1	0	0	0	
21574	1	0	0	0	
21584	1	0	0	0	

	condition_Very Good
5	0
10	0

```

21      0
41      0
70      0
...
21545    0
21552    0
21560    0
21574    0
21584    0

```


```
[1521 rows x 38 columns]
```

We have outliers in 'price', 'sqft_lot', 'sqft_lot15'.

```

In [34]: ▶ # visualizing price ditribution
plt.figure(figsize=(15, 8))
plt.hist(data['price'], bins= 30, color='skyblue', edgecolor='black')
plt.title('Price Distribution')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()

```



The outliers in price are important since they are variations in price levels. For 'sqft_lot', 'sqft_lot15' we may need to perform some transformations on them.

In [35]: ▶ data.describe()

Out[35]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.108200e+04	21082.000000	21082.000000	21082.000000	2.108200e+04	21082.000000
mean	5.402469e+05	3.372403	2.115916	2080.359975	1.507759e+04	1.49362
std	3.667323e+05	0.924996	0.768142	917.856396	4.117338e+04	0.53937
min	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.00000
25%	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.00000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.620000e+03	1.50000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.069775e+04	2.00000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.50000

8 rows × 38 columns



Checking correlations and dealing with multicollinearity

```
In [36]: ▶ # Correlation matrix to see our variable correlations  
correlation_matrix = data.corr()  
correlation_matrix
```


Out[36]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sq
price	1.000000	0.308454	0.525029	0.702004	0.088400	0.256603	(
bedrooms	0.308454	1.000000	0.513694	0.577696	0.032531	0.178518	(
bathrooms	0.525029	0.513694	1.000000	0.754793	0.088451	0.503796	(
sqft_living	0.702004	0.577696	0.754793	1.000000	0.173266	0.354260	(
sqft_lot	0.088400	0.032531	0.088451	0.173266	1.000000	-0.007745	(
floors	0.256603	0.178518	0.503796	0.354260	-0.007745	1.000000	(
sqft_above	0.605481	0.478967	0.685959	0.876787	0.183653	0.523594	-
sqft_basement	0.323018	0.301987	0.281813	0.433369	0.015612	-0.245628	-)
yr_built	0.054849	0.156820	0.508866	0.319584	0.052469	0.489898	(
zipcode	-0.053429	-0.152539	-0.204016	-0.198987	-0.129626	-0.058443	-)
lat	0.307667	-0.009939	0.025243	0.053213	-0.085076	0.049237	-)
long	0.022512	0.131398	0.224660	0.241473	0.230489	0.125360	(
sqft_living15	0.586495	0.391936	0.569396	0.756199	0.143815	0.279379	(
sqft_lot15	0.083530	0.030779	0.089414	0.184920	0.719499	-0.011632	(
grade_10 Very Good	0.341166	0.134985	0.272396	0.368610	0.075398	0.174422	(
grade_11 Excellent	0.356823	0.115891	0.245449	0.344909	0.071959	0.118923	(
grade_12 Luxury	0.287253	0.061427	0.159044	0.238206	0.063029	0.054646	(
grade_13 Mansion	0.214754	0.039577	0.096376	0.146217	0.007920	0.021550	(
grade_3 Poor	-0.005226	-0.017665	-0.012248	-0.011709	-0.000351	-0.006303	-)
grade_4 Low	-0.032053	-0.068905	-0.056341	-0.054607	0.000467	-0.030314	-)
grade_5 Fair	-0.084017	-0.113082	-0.139688	-0.126994	0.021867	-0.079997	-)
grade_6 Low Average	-0.209440	-0.238213	-0.366272	-0.312025	-0.018742	-0.229695	-)
grade_7 Average	-0.317149	-0.107280	-0.314312	-0.359828	-0.066982	-0.309271	-)
grade_8 Good	0.005588	0.075834	0.191163	0.072314	-0.024877	0.201113	(
grade_9 Better	0.236420	0.160343	0.265148	0.318511	0.050922	0.244720	(
view_AVERAGE	0.147555	0.045367	0.085841	0.133146	0.039064	0.006396	(
view_EXCELLENT	0.307035	0.036234	0.108054	0.169713	0.019024	0.025156	(
view_FAIR	0.093931	0.022087	0.038901	0.067767	-0.008165	-0.022713	(
view_GOOD	0.183829	0.049832	0.112348	0.158828	0.069025	0.020403	(
view_NONE	-0.359326	-0.080646	-0.176624	-0.270032	-0.066519	-0.015586	-)
waterfront_NO	-0.055680	0.005788	-0.010212	-0.019120	-0.004858	0.000332	-)
waterfront_Unknown	-0.010632	-0.005528	-0.005646	-0.007231	-0.000528	-0.005499	-)
waterfront_YES	0.260777	-0.001578	0.062055	0.103331	0.021216	0.019853	(
condition_Average	0.009548	0.007366	0.193346	0.105459	-0.011576	0.318246	(

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sq
condition_Fair	-0.052401	-0.049792	-0.076150	-0.064201	0.039403	-0.055165	-0.055165
condition_Good	-0.033639	-0.011579	-0.169355	-0.087109	0.012719	-0.258017	-0.258017
condition_Poor	-0.020132	-0.037211	-0.044078	-0.035674	0.006813	-0.024924	-0.024924
condition_Very Good	0.057935	0.027225	-0.034867	-0.018609	-0.014117	-0.120716	-0.120716

38 rows × 38 columns

```
In [37]: # visualizing the correlations using heatmap
plt.figure(figsize=(30,25))
sns.set(font_scale=1.2)
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap="YlGnBu")
plt.show()
```



Checking highly correlated pairs

```
In [38]: # checking the highly correlated variables
#getting variables with high correlation, having 0.75 as the threshold
threshold = 0.75

# Finding indices where correlation is greater than the threshold and excluding
row, col = np.where((np.abs(correlation_matrix) > threshold) & (np.abs(correlation_matrix) < -threshold))

# Creating a DataFrame with the pairs of variables and their correlation
high_corr_pairs = pd.DataFrame({
    'First_Variable': correlation_matrix.index[row],
    'Second_variable': correlation_matrix.columns[col],
    'Correlation': correlation_matrix.values[row, col]
})

# Display the pairs with high correlation
high_corr_pairs
```

Out[38]:

	First_Variable	Second_variable	Correlation
0	bathrooms	sqft_living	0.754793
1	sqft_living	bathrooms	0.754793
2	sqft_living	sqft_above	0.876787
3	sqft_living	sqft_living15	0.756199
4	sqft_above	sqft_living	0.876787
5	sqft_living15	sqft_living	0.756199
6	waterfront_NO	waterfront_Unknown	-0.967427
7	waterfront_Unknown	waterfront_NO	-0.967427
8	condition_Average	condition_Good	-0.812130
9	condition_Good	condition_Average	-0.812130

To deal with the multicollinearity, we will drop some values causing the multicollinearity.

```
In [39]: # dropping "bathrooms"
data.drop('bathrooms', axis=1, inplace=True)
```

```
In [40]: # dropping "sqft_living15"
data.drop('sqft_living15', axis=1, inplace=True)
```

```
In [41]: # dropping "waterfront_Unknown"
data.drop('waterfront_Unknown', axis=1, inplace=True)
```

```
In [42]: # dropping "condition_Average"
data.drop('condition_Average', axis=1, inplace=True)
```

```
In [43]: # dropping "condition_Good"  
data.drop('condition_Good', axis=1, inplace=True)
```

```
In [44]: # dropping "sqft_lot15" which had outlier  
data.drop('sqft_lot15', axis=1, inplace=True)
```

```
In [45]: # Checking correlations with price  
corr_with_price=data.corr()['price']  
corr_with_price
```

```
Out[45]: price                1.000000  
bedrooms                0.308454  
sqft_living             0.702004  
sqft_lot                0.088400  
floors                 0.256603  
sqft_above             0.605481  
sqft_basement          0.323018  
yr_built               0.054849  
zipcode               -0.053429  
lat                   0.307667  
long                  0.022512  
grade_10 Very Good    0.341166  
grade_11 Excellent    0.356823  
grade_12 Luxury       0.287253  
grade_13 Mansion      0.214754  
grade_3 Poor          -0.005226  
grade_4 Low           -0.032053  
grade_5 Fair          -0.084017  
grade_6 Low Average   -0.209440  
grade_7 Average       -0.317149  
grade_8 Good          0.005588  
grade_9 Better        0.236420  
view_AVERAGE         0.147555  
view_EXCELLENT        0.307035  
view_FAIR             0.093931  
view_GOOD             0.183829  
view_NONE             -0.359326  
waterfront_NO         -0.055680  
waterfront_YES        0.260777  
condition_Fair        -0.052401  
condition_Poor        -0.020132  
condition_Very Good   0.057935  
Name: price, dtype: float64
```

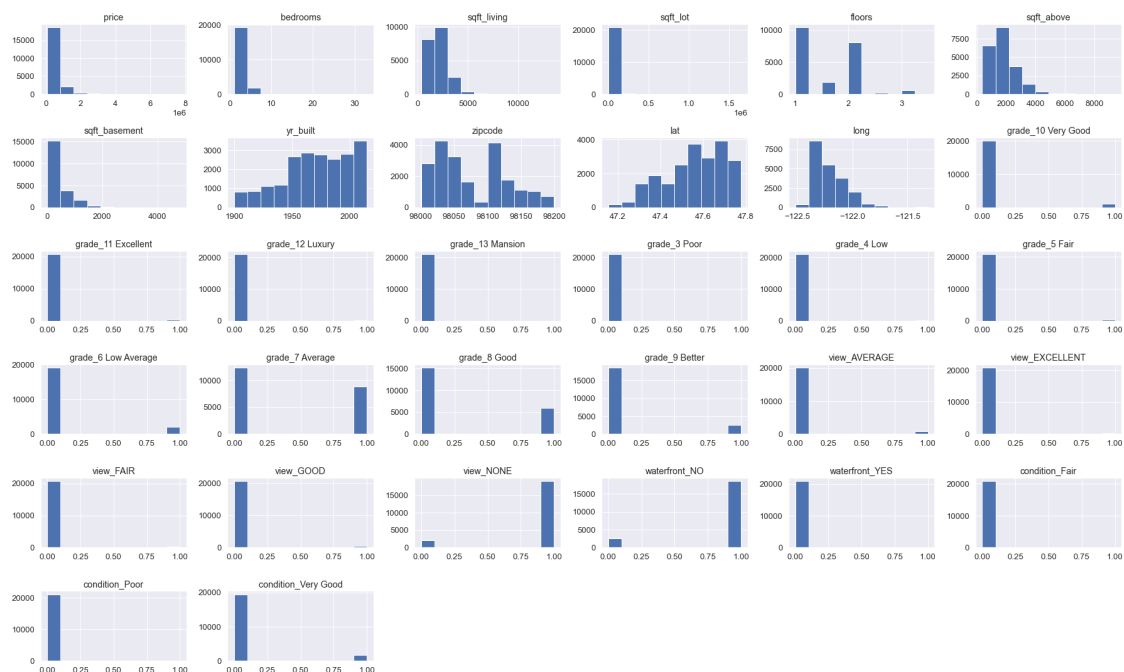


```
In [46]: # plotting correlations with price
plt.figure(figsize=(15, 8))
corr_with_price.drop('price').sort_values().plot(kind='barh')
plt.title('Correlations with Price')
plt.xlabel('Corr Coefficient')
plt.ylabel('Variables')
plt.show();
```



Checking if the data distributions are normal

```
In [47]: # histogram plot for distributions
data.hist(figsize=(25,15))
plt.tight_layout()
plt.show()
```



Most variables dont follow a normal ditribution.

Building Linear Regression Model

Model Iterations

Building a baseline model(model1)

We will use simple linear regression as the baseline model.

```
In [48]: ▶ # importing necessary Libraries  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
import statsmodels.api as sm  
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [49]: data.corr()['price']
```

```
Out[49]: price                1.000000
bedrooms                0.308454
sqft_living             0.702004
sqft_lot                0.088400
floors                  0.256603
sqft_above              0.605481
sqft_basement           0.323018
yr_built                0.054849
zipcode                -0.053429
lat                    0.307667
long                   0.022512
grade_10 Very Good      0.341166
grade_11 Excellent      0.356823
grade_12 Luxury         0.287253
grade_13 Mansion        0.214754
grade_3 Poor            -0.005226
grade_4 Low             -0.032053
grade_5 Fair            -0.084017
grade_6 Low Average     -0.209440
grade_7 Average         -0.317149
grade_8 Good            0.005588
grade_9 Better          0.236420
view_AVERAGE           0.147555
view_EXCELLENT          0.307035
view_FAIR               0.093931
view_GOOD               0.183829
view_NONE               -0.359326
waterfront_NO           -0.055680
waterfront_YES          0.260777
condition_Fair          -0.052401
condition_Poor          -0.020132
condition_Very Good     0.057935
Name: price, dtype: float64
```

For our baseline model we will use the feature 'sqft_living' since it is the most highly correlated with price.

```
In [50]: ▶ # Selecting the dependent and independent variable
X_baseline = data[['sqft_living']]
y = data['price']
# adding a constant for the intercept
baseline_model = sm.OLS(y, sm.add_constant(X_baseline))
#fit the model
baseline_results = baseline_model.fit()
#make predictions
y_pred_baseline = baseline_results.predict(sm.add_constant(X_baseline))
# calculate rmse
baseline_rmse = np.sqrt(mean_squared_error(y, y_pred_baseline))
# displaying results

print(baseline_results.summary())

print(" RMSE for the baseline model:", baseline_rmse)
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.493
Model:                  OLS      Adj. R-squared:
0.493
Method:                 Least Squares    F-statistic:          2.0
48e+04
Date:                   Tue, 02 Jan 2024    Prob (F-statistic):
0.00
Time:                   20:34:32    Log-Likelihood:          -2.92
87e+05
No. Observations:      21082    AIC:          5.8
57e+05
Df Residuals:          21080    BIC:          5.8
58e+05
Df Model:               1
Covariance Type:       nonrobust
=====
=====
               coef      std err          t      P>|t|      [0.025
0.975]
-----
const      -4.327e+04    4456.393     -9.709     0.000    -5.2e+04    -
3.45e+04
sqft_living    280.4877      1.960    143.116     0.000     276.646
284.329
=====
=====
Omnibus:          14303.984    Durbin-Watson:
1.986
Prob(Omnibus):    0.000    Jarque-Bera (JB):          5097
67.330
Skew:             2.786    Prob(JB):
0.00
Kurtosis:         26.437    Cond. No.          5.
63e+03
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.63e+03. This might indicate that there are

strong multicollinearity or other numerical problems.

RMSE for the baseline model: 261170.8023960749

From the first model we note that the R squared is 0.493 to mean that 49.3% of variations in price are explained by square foot living.


The F statistic is 0.00 indicating that the overall model is significant.

The Model RMSE is 261170.8023960749.

We had earlier noted that most variables did not follow a normal distribution 'price' being one of them. We will therefore log transform price to see if the model improves.

Model 2

Here we are inspecting how the model performs with only the 'price' transformed.

```
In [51]:  # Selecting the dependent and independent variable  
X_baseline = data[['sqft_living']]  
y = np.log(data['price']+1)  
# adding a constant for the intercept  
baseline_model = sm.OLS(y, sm.add_constant(X_baseline))  
#fit the model  
baseline_results = baseline_model.fit()  
#make predictions  
y_pred_baseline =baseline_results.predict(sm.add_constant(X_baseline))  
# calculate rmse  
baseline_rmse = np.sqrt(mean_squared_error(y, y_pred_baseline))  
# displaying results  
  
print(baseline_results.summary())  
  
print(" RMSE for the baseline model:", baseline_rmse)
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.483
Model:                  OLS      Adj. R-squared:
0.483
Method:                 Least Squares    F-statistic:          1.9
70e+04
Date:                   Tue, 02 Jan 2024    Prob (F-statistic):
0.00
Time:                   20:34:32    Log-Likelihood:          -
9429.6
No. Observations:      21082    AIC:          1.8
86e+04
Df Residuals:          21080    BIC:          1.8
88e+04
Df Model:               1
Covariance Type:       nonrobust
=====
=====
               coef      std err          t      P>|t|      [0.025
0.975]
-----
const          12.2190      0.006    1892.178      0.000      12.206
12.232
sqft_living     0.0004    2.84e-06    140.355      0.000      0.000
0.000
=====
=====
Omnibus:          3.289    Durbin-Watson:
1.981
Prob(Omnibus):    0.193    Jarque-Bera (JB):
3.309
Skew:             0.029    Prob(JB):
0.191
Kurtosis:         2.982    Cond. No.          5.
63e+03
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.63e+03. This might indicate that there are


strong multicollinearity or other numerical problems.

RMSE for the baseline model: 0.3784548319492928

The square foot of living now explains 48.3% (R squared) of variations in price. We also still have an error 'The condition number is large, 5.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.' We will then explore how the model performs after transforming both the feature and target variable.

Model 3

Here we have both 'sqft_living ' and 'price transformed'

```
In [52]:  # Selecting the dependent and independent variable  
X = np.log(data[['sqft_living']])  
y = np.log(data['price']+1)  
# adding a constant for the intercept  
model = sm.OLS(y, sm.add_constant(X))  
#fit the model  
results = model.fit()  
#make predictions  
y_pred = results.predict(sm.add_constant(X))  
# calculate rmse  
rmse = np.sqrt(mean_squared_error(y, y_pred))  
# displaying results  
  
print(results.summary())  
  
print(" RMSE for the baseline model:", rmse)
```

OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.455
Model:                  OLS      Adj. R-squared:
0.455
Method:                 Least Squares    F-statistic:          1.7
59e+04
Date:                   Tue, 02 Jan 2024    Prob (F-statistic):
0.00
Time:                   20:34:32    Log-Likelihood:          -
9989.4
No. Observations:      21082    AIC:          1.9
98e+04
Df Residuals:          21080    BIC:          2.0
00e+04
Df Model:              1
Covariance Type:       nonrobust
=====
=====
               coef      std err          t      P>|t|      [0.025
0.975]
-----
const          6.7255      0.048    140.854      0.000      6.632
6.819
sqft_living    0.8374      0.006    132.627      0.000      0.825
0.850
=====
=====
Omnibus:          121.179    Durbin-Watson:
1.980
Prob(Omnibus):    0.000    Jarque-Bera (JB):          1
12.125
Skew:            0.144    Prob(JB):          4.
49e-25
Kurtosis:        2.789    Cond. No.
137.
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

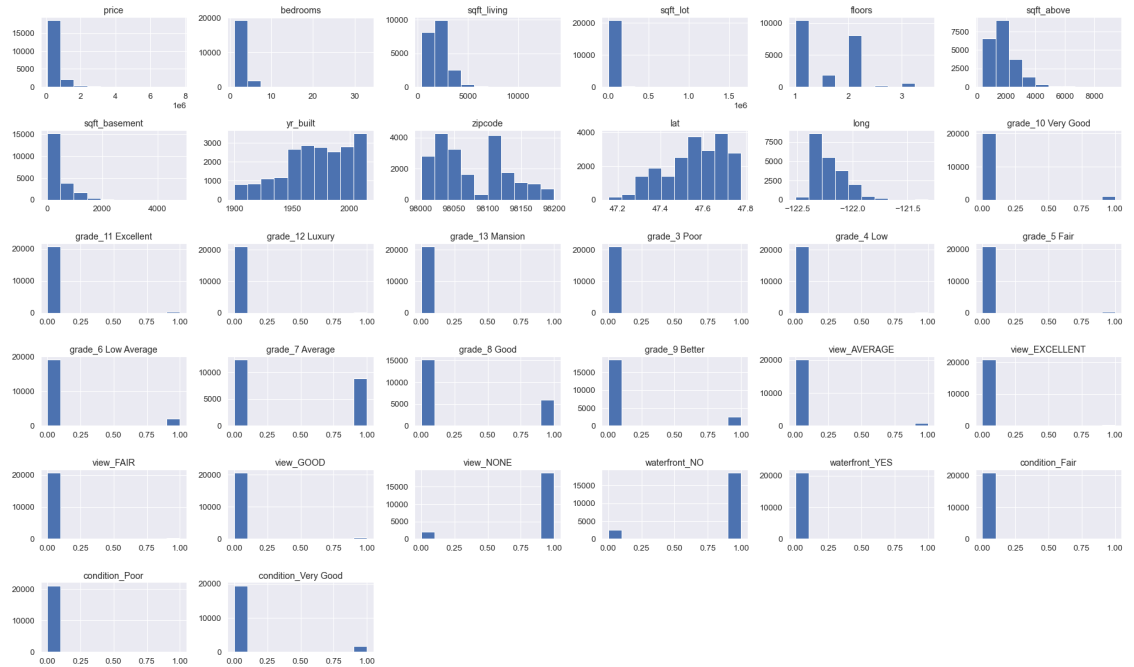
RMSE for the baseline model: 0.3886403105841183

For the transformed variables, the target variable(price) is now explained by 45.5%(R squared) in price. We also note that the error we were getting that (there is a possibility of strong multicollinearity or other numeric problems) has been resolved.

In the next model we will try transform multiple features that do not follow a normal distribution and add them to our model. Then inspect how our model performs.


Before log transformation

```
In [53]: ▶ # histogram plot for distributions  
data.hist(figsize=(25,15))  
plt.tight_layout()  
plt.show()
```



After log transformation

```
In [54]: ▶ # Log transformation to normalize the variables and rename them
data["log_price"] = np.log(data["price"]+1)
data["log_sqft_living"] = np.log(data["sqft_living"]+1)
data["log_sqft_lot"] = np.log(data["sqft_lot"]+1)
data["log_sqft_above"] = np.log(data["sqft_above"]+1)
data["log_waterfront_YES"] = np.log(data["waterfront_YES"]+1)
data["log_floors"] = np.log(data["floors"]+1)
# checking the transformed
plot_data = data[["log_price", 'log_sqft_living', 'log_sqft_lot', 'log_sqft_a
plot_data.hist(figsize=(25,15))
plt.tight_layout()
plt.show()
```



Model 4

```
In [55]: ▶ # Selecting independent and dependent variables and using some transforms
X = data[['log_sqft_living', 'waterfront_YES', 'view_EXCELLENT', 'condition_10 Very Good', 'grade_9 Better', 'grade_10 Very Good', 'grade_11 Excellent', 'grade_12 Excellent', 'grade_13 Mansion', 'log_sqft_above', 'log_sqft_lot']]

y = data['log_price']

# Adding a constant term for the intercept in the multiple regression model
model=sm.OLS(y, sm.add_constant(X))

# Fitting the multiple regression model
results = model.fit()
#making predictions
y_pred=results.predict(sm.add_constant(X))
#calculating rsme
rmse=np.sqrt(mean_squared_error(y, y_pred))

# Display the summary of the regression and rmse
print(results.summary())
print(" RMSE for the baseline model:", rmse)
```

OLS Regression Results

```

=====
=====
Dep. Variable:          log_price    R-squared:
0.561
Model:                  OLS         Adj. R-squared:
0.561
Method:                 Least Squares    F-statistic:
2243.
Date:                   Tue, 02 Jan 2024    Prob (F-statistic):
0.00
Time:                   20:34:42    Log-Likelihood:        -
7708.2
No. Observations:      21082    AIC:                        1.5
44e+04
Df Residuals:          21069    BIC:                        1.5
55e+04
Df Model:               12
Covariance Type:        nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.
025      0.975]
-----
const                9.1349      0.058    157.130    0.000      9.
021      9.249
log_sqft_living      0.7078      0.012     60.921    0.000      0.
685      0.731
waterfront_YES      0.4086      0.036     11.420    0.000      0.
338      0.479
view_EXCELLENT      0.2958      0.025     12.066    0.000      0.
248      0.344
condition_Very Good  0.1580      0.009     17.512    0.000      0.
140      0.176
grade_7 Average     -0.0812      0.005    -14.891    0.000     -0.
092     -0.071
grade_9 Better       0.2741      0.009     31.429    0.000      0.
257      0.291
grade_10 Very Good   0.4745      0.012     38.355    0.000      0.
450      0.499
grade_11 Excellent   0.6763      0.019     34.757    0.000      0.
638      0.714
grade_12 Luxury      0.8849      0.039     22.912    0.000      0.
809      0.961
grade_13 Mansion     1.2596      0.098     12.909    0.000      1.
068      1.451
log_sqft_above      -0.1240      0.012    -10.401    0.000     -0.
147     -0.101
log_sqft_lot        -0.0640      0.003    -22.503    0.000     -0.
070     -0.058
=====
=====
Omnibus:              10.330    Durbin-Watson:
1.976
Prob(Omnibus):         0.006    Jarque-Bera (JB):
10.016

```

```
Skew:                0.037    Prob(JB):
0.00668
Kurtosis:            2.923    Cond. No.
568.
```

```
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

RMSE for the baseline model: 0.34878164210142815

After transforming and adding more features, R squared and adjusted R squared have now increased to 56.1%. Meaning that 56.1% of variations in price are now explained by the independent variables. The F statistic probability is 0.00 to mean that the model overall is significant. RMSE is also now at 0.34878164210142815 which is less than what we had in the log transformed baseline model which we found rmse as 0.3886403105841183. This means that our model accuracy has improved.


Checking Regression Assumptions

We are going to check if the Regression model has passed the assumptions before doing interpretation of the results.

We will inspect **Linearity, Independence, Normality and Equal Variance**

Linearity

```
In [56]: ▶ # plotting model results
fig, ax=plt.subplots()
ax.scatter(y, results.resid, color='green')
ax.axhline(y=0, color='black')
ax.set_xlabel('y')
ax.set_ylabel('residuals')
ax.set_title('Linearity Residual plot');
```



The points form a curvature to mean that the linearity assumption is met

Rainbow stat-test for linearity

```
In [57]: ▶ # performing a rainbow test to test linearity statistically
from statsmodels.stats.diagnostic import linear_rainbow
linear_rainbow(results)
```

```
Out[57]: (0.9485833390658518, 0.9966225779067938)
```


The p value is close to 1. This high p-value indicates that there is not enough evidence to reject the null hypothesis of linearity. Therefore, based on this test, the assumption of linearity is considered to be met.

Independence

The Durbin-Watson statistic is around 1.976 which suggests little to no autocorrelation in the residuals.

The Normality Assumption

```
In [58]: ▶ import scipy.stats as stats
residuals = results.resid
fig = sm.graphics.qqplot(residuals, dist=stats.norm, line='45', fit=True)
fig.suptitle('Residuals QQ Plot')
fig.set_size_inches(10, 5)
plt.show()
```




From the Q-Q plot, we see that the residuals follow a normal distribution. We can conclude that normality assumption is considered met.

The Homoscedasticity Assumption(Equal Variance)

```
In [59]: # scatter plot to check homoscedasticity
plt.figure(figsize=(10, 6))
sns.scatterplot(x=data['log_price'], y=results.resid)
plt.title('Residuals vs. Predicted Values')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
```

```
Out[59]: Text(0, 0.5, 'Residuals')
```



From the scatter plot we observe that there is little to no heteroscedasticity in the residuals.

Interpretation of results

Baseline Model: R-squared: 0.493 Adjusted R-squared: 0.493 RMSE: 261170.80

Model 2 (log-transformed price): R-squared: 0.483 Adjusted R-squared: 0.483 RMSE: 0.3785

Model 3 (log-transformed price and sqft_living): R-squared: 0.455 Adjusted R-squared: 0.455 RMSE: 0.3886

Model 4 (multiple features and log-transformed price): R-squared: 0.561 Adjusted R-squared: 0.561 RMSE: 0.3488

Analysis Interpretation: The R-squared values provide a measure of how well the models explain the variations in the target variable (price). As we progress from the baseline to the 4th model, the R-squared increases, indicating better explanatory power.

The RMSE values for the log-transformed models (Model 2 and Model 3), the RMSE is significantly lower than the baseline, indicating better predictive performance.

Model 4, which includes multiple features, the R-squared further improves, and the RMSE decreases compared to the log-transformed models. This suggests that the inclusion of additional features has enhanced the model's ability to predict prices.

Interpretation:

Model 4 with multiple features and log-transformed price performs better than the baseline model, both in terms of explanatory power and predictive accuracy. The probability F statistic being 0.00 means that the model overall is significant. The P values for our coefficients all being 0.00 means that the coefficients as well are significant for our test.

Interpreting coefficients

grade_13 Mansion (Coefficient: 1.2596): A one-unit increase in the presence of the "Mansion" grade is associated with an estimated increase of approximately 1.2596 units in the log of house prices. This variable has the highest positive coefficient.

grade_12 Luxury (Coefficient: 0.8849): one-unit increase in the presence of the "Luxury" grade is associated with an estimated increase of approximately 0.8849 units in the log of house prices. The "Luxury" grade has the second-highest positive coefficient.

grade_11 Excellent (Coefficient: 0.6763): A one-unit increase in the presence of the "Excellent" grade is associated with an estimated increase of approximately 0.6763 units in the log of house prices. Houses with an "Excellent" grade have the third-highest positive coefficient.

log_sqft_living (Coefficient: 0.7078): A one-unit increase in the logarithm of square footage living area is associated with an estimated increase of approximately 0.7078 units in the log of house prices. The logarithm of square footage living area has a positive impact.

grade_10 Very Good (Coefficient: 0.4745): A one-unit increase in the presence of the "Very Good" grade is associated with an estimated increase of approximately 0.4745 units in the log of house prices. Houses with a "Very Good" grade contribute positively.

view_EXCELLENT (Coefficient: 0.2958): A one-unit increase in the presence of an "Excellent" view is associated with an estimated increase of approximately 0.2958 units in the log of house prices. Houses with an "Excellent" view contribute positively.

waterfront_YES (Coefficient: 0.4086): A one-unit increase in the presence of a waterfront is associated with an estimated increase of approximately 0.4086 units in the log of house prices. Houses with a waterfront contribute positively.

grade_9 Better (Coefficient: 0.2741): A one-unit increase in the presence of the "Better" grade is associated with an estimated increase of approximately 0.2741 units in the log of house prices. Houses with a "Better" grade contribute positively.

condition_Very Good (Coefficient: 0.1580): A one-unit increase in the presence of a "Very Good" condition is associated with an estimated increase of approximately 0.1580 units in the log of house prices. Houses in very good condition contribute positively.

log_sqft_above (Coefficient: -0.1240): A one-unit increase in the logarithm of square footage above is associated with an estimated decrease of approximately 0.1240 units in the log of house prices. The logarithm of square footage of the lot above has a negative impact.

grade_7 Average (Coefficient: -0.0812): A one-unit increase in the presence of the "Average" grade is associated with an estimated decrease of approximately 0.0812 units in the log of house prices. Houses with an "Average" grade (grade 7) contribute negatively.

log_sqft_lot (Coefficient: -0.0640): A one-unit increase in the logarithm of square footage of the lot is associated with an estimated decrease of approximately 0.0640 units in the log of house prices. The logarithm of square footage of the lot has a negative impact.

Summary

The features associated with higher-grade classifications (grade_13 Mansion, grade_11 Excellent, grade_12 Luxury) and larger living area (log_sqft_living) have the most positive impact on house prices, while features like lower-grade classifications (grade_7 Average) and smaller square footage above ground (log_sqft_above) have a negative impact.

Answering objectives

What are the key features that influence house prices

The features associated with higher-grade classifications (grade_13 Mansion, grade_11 Excellent, grade_12 Luxury) and larger living area (log_sqft_living) have the most positive impact on house prices, while features like lower-grade classifications (grade_7 Average) and smaller square footage above ground (log_sqft_above) have a negative impact.

What Feature has the highest impact on house prices

Houses with a grade_13 Mansion (Coefficient: 1.2596) had the highest influence of house prices.

Evaluating and validating the performance of the model.

The study developed multiple predictive models with increasing complexity, including additional log-transformed features and log-transformed price. The models were evaluated using metrics such as R-squared and RMSE to assess their explanatory power and predictive accuracy. The improvement in R-squared values and the reduction in RMSE indicate successful model development and validation.

Recommendations from our study

- Grade has been identified to have the most impact on House prices. This includes various factors such as the quality of construction, materials used, architectural design, and overall condition. Real estate investors seeking premium returns should consider the grade of the house.
- Real estate investors should also consider waterfront locations and excellent views as they also impact prices.
- Real estate investors should recognize the positive impact of larger living areas, as indicated by the `log_sqft_living` variable in order to fetch higher returns.
- Investors should be mindful of features with a negative impact on house prices, such as lower-grade classifications ("Average") and smaller square footage above ground (`log_sqft_above`).

Limitations of the study

- The study does not consider external factors such as economic policies, interest rates, or global economic conditions, which can influence the real estate market.
- While the analysis identifies associations between features and house prices, it does not establish causation. The observed relationships may be influenced by confounding factors not included in the model
- The analysis assumes a linear relationship between the independent variables and the house prices. Non-linear relationships or interactions between variables might not be fully captured.
- Linear regression assumes continuous independent variables. While categorical variables can be included using dummy coding, this approach might not capture the full complexity of categorical relationships.

Steps to consider based on Limitations

- Consider integrating macro economic data and other external factors that affect house prices
- Consider employing non-linear regression models or machine learning algorithms that can capture non-linear relationships between variables.